INSTALLATION AND CONFIGURATION

Installation and Configuration

Installing kubectl

To configure and manage your cluster, you will probably use the **kubectl** command. You can use RESTful calls or the Go language, as well.

Enterprise Linux distributions have the various Kubernetes utilities and other files available in their repositories. For example, on RHEL/CentOS, you would find **kubectl** in the **kubernetes-client** package. On OpenShift they use a command very similar to **kubectl** called **oc**.

You can (if needed) download the code from GitHub, and go through the usual steps to compile and install **kubectl**.

This command line will use **$HOME/.kube/config** as a configuration file. This contains all the Kubernetes endpoints that you might use. If you examine it, you will see cluster definitions (i.e. IP endpoints), credentials, and contexts.

A *context* is a combination of a cluster and user credentials. You can pass these parameters on the command line, or switch the shell between contexts with a command, as in:

```
$ kubectl config use-context foobar
```

This is handy when going from a local environment to a cluster in the cloud, or from one cluster to another, such as from development to production.