



SCHEDULING

Scheduling

Filtering (Predicates)

The scheduler goes through a set of filters, or predicates, to find available nodes, then ranks each node using priority functions. The node with the highest rank is selected to run the Pod.

```
predicatesOrdering = []string{CheckNodeConditionPred, GeneralPred,  
HostNamePred, PodFitsHostPortsPred, MatchNodeSelectorPred, PodFitsResourcesPred,  
NoDiskConflictPred, PodToleratesNodeTaintsPred,  
PodToleratesNodeNoExecuteTaintsPred, CheckNodeLabelPresencePred,  
checkServiceAffinityPred, MaxEBSVolumeCountPred,  
MaxGCEPDVolumeCountPred, MaxAzureDiskVolumeCountPred,  
CheckVolumeBindingPred, NoVolumeZoneConflictPred,  
CheckNodeMemoryPressurePred, CheckNodeDiskPressurePred,  
MatchInterPodAffinityPred}
```

The predicates, such as **PodFitsHost** or **NoDiskConflict**, are evaluated in a particular and configurable order. In this way, a node has the least amount of checks for new Pod deployment, which can be useful to exclude a node from unnecessary checks if the node is not in the proper condition.

For example, there is a filter called **HostNamePred**, which is also known as **HostName**, which filters out nodes that do not match the node name specified in the pod specification. Another predicate is **PodFitsResources** to make sure that the available CPU and memory can fit the resources required by the Pod.

The scheduler can be updated by passing a configuration of **kind: Policy**, which can order predicates, give special weights to priorities, and even **hardPodAffinitySymmetricWeight**, which deploys Pods such that if we set Pod A to run with Pod B, then Pod B should automatically be run with Pod A.