



## SCHEDULING

## Scheduling

### Pod Affinity Rules

Pods which may communicate a lot or share data may operate best if co-located, which would be a form of affinity. For greater fault tolerance, you may want Pods to be as separate as possible, which would be anti-affinity. These settings are used by the scheduler based on the labels of Pods that are already running. As a result, the scheduler must interrogate each node and track the labels of running Pods. Clusters larger than several hundred nodes may see significant performance loss. Pod affinity rules use **In**, **NotIn**, **Exists**, and **DoesNotExist** operators.

*Click on the boxes to learn more about Pod affinity rules*

#### Pod Affinity Rules

Close ^

##### `requiredDuringSchedulingIgnoredDuringExecution`

The use of **`requiredDuringSchedulingIgnoredDuringExecution`** means that the Pod will not be scheduled on a node unless the following operator is true. If the operator changes to become false in the future, the Pod will continue to run. This could be seen as a hard rule.

Close ^

##### `preferredDuringSchedulingIgnoredDuringExecution`

Similarly, **`preferredDuringSchedulingIgnoredDuringExecution`** will choose a node with the desired setting before those without. If no properly-labeled nodes are available, the Pod will execute anyway. This is more of a soft setting, which declares a preference instead of a requirement.

Close ^

##### `podAffinity`

With the use of **podAffinity**, the scheduler will try to schedule Pods together.

### podAntiAffinity

Close ^

The use of **podAntiAffinity** would cause the scheduler to keep Pods on different nodes.