SCHEDULING

Scheduling

Node Affinity Rules

Where Pod affinity/anti-affinity has to do with other Pods, the use of **`nodeAffinity`** allows Pod scheduling based on node labels. This is similar and will some day replace the use of the **`nodeSelector`** setting. The scheduler will not look at other Pods on the system, but the labels of the nodes. This should have much less performance impact on the cluster, even with a large number of nodes.

- Uses **`In`**, **`NotIn`**, **`Exists`**, **`DoesNotExist`** operators
- **`requiredDuringSchedulingIgnoredDuringExecution`**
- **`preferredDuringSchedulingIgnoredDuringExecution`**
- Planned for future: **`requiredDuringSchedulingRequiredDuringExecution`**.

Until **`nodeSelector`** has been fully deprecated, both the selector and required labels must be met for a Pod to be scheduled.