



LOGGING AND TROUBLESHOOTING

Logging and Troubleshooting

Cluster Start Sequence

The cluster startup sequence begins with systemd if you built the cluster using **kubeadm**. Other tools may leverage a different method. Use **systemctl status kubelet.service** to see the current state and configuration files used to run the kubelet binary.

- Uses **/etc/systemd/system/kubelet.service.d/10-kubeadm.conf** file

Inside of the **config.yaml** file you will find several settings for the binary, including the **staticPodPath** which indicates the directory where kubelet will read every yaml file and start every pod. If you put a yaml file in this directory, it is a way to troubleshoot the scheduler, as the pod is created with any requests to the scheduler.

- Uses **/var/lib/kubelet/config.yaml** configuration file
- **staticPodPath** is set to **/etc/kubernetes/manifests/**

The four default yaml files will start the base pods necessary to run the cluster:

- kubelet creates all pods from *.yaml in directory: kube-apiserver, etcd, kube-controller-manager, kube-scheduler.

Once the watch loops and controllers from kube-controller-manager run using etcd data, the rest of the configured objects will be created.